

```
/home/zsakkaff/ai-science-training-  
series/07_largeScaleTraining/src/ai4sci/logs/2022-11-01-154104-  
thetagpu06_ngpu1_ncpu256.log  
/home/zsakkaff/ai-science-training-  
series/07_largeScaleTraining/src/ai4sci/logs/2022-11-01-184048-  
thetagpu06_ngpu1_ncpu256.log  
/home/zsakkaff/ai-science-training-  
series/07_largeScaleTraining/src/ai4sci/logs/2022-11-01-184210-  
thetagpu06_ngpu1_ncpu256.log  
/home/zsakkaff/ai-science-training-  
series/07_largeScaleTraining/src/ai4sci/logs/2022-11-01-225933-  
thetagpu06_ngpu1_ncpu256.log
```

```
zsakkaiff@thetagpu06:~/ai-science-training-series/07_largeScaleTraining/src/ai4sci$ ./main.sh batch_size=512 epochs=5 > main.log 2>&1 &
[1] 1526218
zsakkaiff@thetagpu06:~/ai-science-training-series/07_largeScaleTraining/src/ai4sci$ tail -f main.log $(tail -1 logs/latest)
==> main.log <==
- NGPUS TOTAL: 1
- MAIN: /home/zsakkaiff/ai-science-training-series/07_largeScaleTraining/src/ai4sci/main.py
- Writing logs to /home/zsakkaiff/ai-science-training-series/07_largeScaleTraining/src/ai4sci/logs/2022-11-01-225933-thet
agpu06_ngpu1_ncpu256.log
- python3: /lus/theta-fs0/software/thetagpu/conda/2022-07-01/mconda3/bin/python3
- mpirun: /lus/theta-fs0/software/thetagpu/openmpi/openmpi-4.1.4_ucx-1.12.1_gcc-9.4.0/bin/mpirun
- ai4sci: /home/zsakkaiff/ai-science-training-series/07_largeScaleTraining/src/ai4sci/__init__.py
- exec: /lus/theta-fs0/software/thetagpu/openmpi/openmpi-4.1.4_ucx-1.12.1_gcc-9.4.0/bin/mpirun --verbose -n 1 --
hostfile /var/tmp/cobalt.10112460 -npernode 1 -x PYTHONUSERBASE -x PYTHONSTARTUP -x http_proxy -x https
_proxy -x PATH -x LD_LIBRARY_PATH /lus/theta-fs0/software/thetagpu/conda/2022-07-01/mconda3/bin/python3 /home/zsakk
aiff/ai-science-training-series/07_largeScaleTraining/src/ai4sci/main.py batch_size=512 epochs=5
```

```
2022-11-01 23:00:01.564837: W tensorflow/core/grappler/costs/op_level_cost_estimator.cc:690] Error in PredictCost() for the
op: op: "CropAndResize" attr { key: "T" value { type: DT_UINT8 } } attr { key: "extrapolation_value" value { f: 0 } } attr
{ key: "method" value { s: "bilinear" } } inputs { dtype: DT_UINT8 shape { dim { size: 1 } dim { size: -5 } dim { size: -6
} dim { size: 3 } } } inputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: 4 } } } inputs { dtype: DT_INT32 shape
{ dim { size: -2 } } } inputs { dtype: DT_INT32 shape { dim { size: 2 } } } device { type: "CPU" vendor: "AuthenticAMD" mo
del: "241" frequency: 2245 num_cores: 2 environment { key: "cpu_instruction_set" value: "SSE, SSE2, SSE3" } environment { k
ey: "eigen" value: "3.4.90" } l1_cache_size: 32768 l2_cache_size: 524288 l3_cache_size: 268435456 memory_size: 268435456 }
outputs { dtype: DT_FLOAT shape { dim { size: -2 } dim { size: -10 } dim { size: -11 } dim { size: 3 } } }
```

```
==> /home/zsakkaiff/ai-science-training-series/07_largeScaleTraining/src/ai4sci/logs/2022-11-01-225933-thetagpu06_ngpu1_ncpu
256.log <==
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 29/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 30/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 31/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Trees [0] -1/-1/-1->0->-1 [1] -1/-1/-1->0->-1 [2] -1/-1/-1->0->-1 [3] -1/-1/-1->0-
>-1 [4] -1/-1/-1->0->-1 [5] -1/-1/-1->0->-1 [6] -1/-1/-1->0->-1 [7] -1/-1/-1->0->-1 [8] -1/-1/-1->0->-1 [9] -1/-1/-1->0->-1
[10] -1/-1/-1->0->-1 [11] -1/-1/-1->0->-1 [12] -1/-1/-1->0->-1 [13] -1/-1/-1->0->-1 [14] -1/-1/-1->0->-1 [15] -1/-1/-1->0-
>-1 [16] -1/-1/-1->0->-1 [17] -1/-1/-1->0->-1 [18] -1/-1/-1->0->-1 [19] -1/-1/-1->0->-1 [20] -1/-1/-1->0->-1 [21] -1/-1/-1-
>0->-1 [22] -1/-1/-1->0->-1 [23] -1/-1/-1->0->-1 [24] -1/-1/-1->0->-1 [25] -1/-1/-1->0->-1 [26] -1/-1/-1->0->-1 [27] -1/-1/-
1->0->-1 [28] -1/-1/-1->0->-1 [29] -1/-1/-1->0->-1 [30] -1/-1/-1->0->-1 [31] -1/-1/-1->0->-1
thetagpu06:1527708:1527835 [0] NCCL INFO Connected all rings
thetagpu06:1527708:1527835 [0] NCCL INFO Connected all trees
thetagpu06:1527708:1527835 [0] NCCL INFO 32 coll channels, 32 p2p channels, 32 p2p channels per peer
thetagpu06:1527708:1527835 [0] NCCL INFO comm 0x7f880aff6d0 rank 0 nranks 1 cudaDev 0 busId f000 - Init COMPLETE
[2022-11-01 23:00:41,917][ai4sci.trainer][INFO] - [0] [0/5: 0 / 1281167 (0%)] epoch=0.0000 dt=37.8666 running_loss=0.0000 b
atch_loss=0.0151 acc=0.0000 batch_acc=0.0039
[2022-11-01 23:00:56,684][ai4sci.trainer][INFO] - [0] [0/5: 5120 / 1281167 (0%)] epoch=0.0000 dt=0.7848 running_loss=0.0000
batch_loss=0.0095 acc=0.0001 batch_acc=0.0293
[2022-11-01 23:01:17,685][ai4sci.trainer][INFO] - [0] [0/5: 10240 / 1281167 (1%)] epoch=0.0000 dt=0.6010 running_loss=0.000
1 batch_loss=0.0095 acc=0.0002 batch_acc=0.0430
[2022-11-01 23:01:41,301][ai4sci.trainer][INFO] - [0] [0/5: 15360 / 1281167 (1%)] epoch=0.0000 dt=0.7152 running_loss=0.000
1 batch_loss=0.0093 acc=0.0004 batch_acc=0.0449
[2022-11-01 23:02:04,130][ai4sci.trainer][INFO] - [0] [0/5: 20480 / 1281167 (2%)] epoch=0.0000 dt=0.7532 running_loss=0.000
2 batch_loss=0.0095 acc=0.0006 batch_acc=0.0410
```

2022-11-01-184048-thetagpu06_ngpu1_ncpu256.log

```
{
  "data": {
    "handler": "ilsvrc_dataset",
    "batch_size": 512,
    "train_filelist":
"/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
    "test_filelist":
"/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
    "shuffle_buffer": 200000,
    "reshuffle_each_iteration": true,
    "num_parallel_readers": 128,
    "prefetch_buffer_size": 8,
    "crop_image_size": [
      256,
      256
    ],
    "num_classes": 1000,
    "num_channels": 3
  }
}
num labels: 1000
build dataset /grand/projects/ALCFAITP/ilsvrc_train_filelist.txt
[36m2022-11-01
18:41:40,382[0m][34mai4sci.ilsvrc_dataset[0m][32mINFO[0m - input
filelist contains 1281167 files, estimated batches per rank 2502[0m
build dataset /grand/projects/ALCFAITP/ilsvrc_val_filelist.txt
[36m2022-11-01
18:41:42,636[0m][34mai4sci.ilsvrc_dataset[0m][32mINFO[0m - input
filelist contains 50000 files, estimated batches per rank 97[0m
thetagpu06:1136611:1136854 [0] NCCL INFO Bootstrap : Using
enp226s0:10.230.2.194<0>
thetagpu06:1136611:1136854 [0] NCCL INFO NET/Plugin : No plugin found
(libnccl-net.so), using internal implementation
thetagpu06:1136611:1136854 [0] NCCL INFO NET/IB : Using [0]mlx5_0:1/IB
[1]mlx5_1:1/IB [2]mlx5_2:1/IB [3]mlx5_3:1/IB [4]mlx5_4:1/IB
[5]mlx5_5:1/RoCE [6]mlx5_6:1/IB [7]mlx5_7:1/IB [8]mlx5_8:1/IB
[9]mlx5_9:1/IB [10]mlx5_10:1/IB [11]mlx5_11:1/RoCE [RO]; OOB
enp226s0:10.230.2.194<0>
thetagpu06:1136611:1136854 [0] NCCL INFO Using network IB
NCCL version 2.12.12+cuda11.0
thetagpu06:1136611:1136854 [0] NCCL INFO Setting affinity for GPU 0 to
010000,00000000,00000000,00000000,00010000,00000000
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 00/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 01/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 02/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 03/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 04/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 05/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 06/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 07/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 08/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 09/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 10/32 : 0
```

```
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 11/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 12/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 13/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 14/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 15/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 16/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 17/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 18/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 19/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 20/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 21/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 22/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 23/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 24/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 25/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 26/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 27/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 28/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 29/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 30/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Channel 31/32 : 0
thetagpu06:1136611:1136854 [0] NCCL INFO Trees [0] -1/-1/-1->0->-1 [1] -
1/-1/-1->0->-1 [2] -1/-1/-1->0->-1 [3] -1/-1/-1->0->-1 [4] -1/-1/-1->0->-
1 [5] -1/-1/-1->0->-1 [6] -1/-1/-1->0->-1 [7] -1/-1/-1->0->-1 [8] -1/-1/-
1->0->-1 [9] -1/-1/-1->0->-1 [10] -1/-1/-1->0->-1 [11] -1/-1/-1->0->-1
[12] -1/-1/-1->0->-1 [13] -1/-1/-1->0->-1 [14] -1/-1/-1->0->-1 [15] -1/-
1/-1->0->-1 [16] -1/-1/-1->0->-1 [17] -1/-1/-1->0->-1 [18] -1/-1/-1->0->-
1 [19] -1/-1/-1->0->-1 [20] -1/-1/-1->0->-1 [21] -1/-1/-1->0->-1 [22] -
1/-1/-1->0->-1 [23] -1/-1/-1->0->-1 [24] -1/-1/-1->0->-1 [25] -1/-1/-1-
>0->-1 [26] -1/-1/-1->0->-1 [27] -1/-1/-1->0->-1 [28] -1/-1/-1->0->-1
[29] -1/-1/-1->0->-1 [30] -1/-1/-1->0->-1 [31] -1/-1/-1->0->-1
thetagpu06:1136611:1136854 [0] NCCL INFO Connected all rings
thetagpu06:1136611:1136854 [0] NCCL INFO Connected all trees
thetagpu06:1136611:1136854 [0] NCCL INFO 32 coll channels, 32 p2p
channels, 32 p2p channels per peer
thetagpu06:1136611:1136854 [0] NCCL INFO comm 0x7f2f9effe6d0 rank 0
nranks 1 cudaDev 0 busId f000 - Init COMPLETE
[36m2022-11-01
18:42:24,498[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1: 0 /
1281167 (0%)] epoch=0.0000 dt=38.0835 running_loss=0.0000
batch_loss=0.0149 acc=0.0000 batch_acc=0.0059[0m
[36m2022-11-01
18:42:49,851[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
5120 / 1281167 (0%)] epoch=0.0000 dt=2.4088 running_loss=0.0000
batch_loss=0.0100 acc=0.0001 batch_acc=0.0176[0m
[36m2022-11-01
18:43:16,391[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
10240 / 1281167 (1%)] epoch=0.0000 dt=0.6993 running_loss=0.0001
batch_loss=0.0092 acc=0.0002 batch_acc=0.0371[0m
[36m2022-11-01
18:43:48,253[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
15360 / 1281167 (1%)] epoch=0.0000 dt=0.8567 running_loss=0.0001
batch_loss=0.0094 acc=0.0004 batch_acc=0.0332[0m
```

```
[36m2022-11-01
18:44:17,265[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
20480 / 1281167 (2%)] epoch=0.0000 dt=0.6943 running_loss=0.0002
batch_loss=0.0092 acc=0.0006 batch_acc=0.0645[0m
[36m2022-11-01
18:44:44,700[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
25600 / 1281167 (2%)] epoch=0.0000 dt=0.6001 running_loss=0.0002
batch_loss=0.0094 acc=0.0008 batch_acc=0.0371[0m
[36m2022-11-01
18:45:11,450[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
30720 / 1281167 (2%)] epoch=0.0000 dt=0.6313 running_loss=0.0002
batch_loss=0.0089 acc=0.0010 batch_acc=0.0703[0m
[36m2022-11-01
18:45:41,162[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
35840 / 1281167 (3%)] epoch=0.0000 dt=0.5692 running_loss=0.0003
batch_loss=0.0089 acc=0.0013 batch_acc=0.0781[0m
[36m2022-11-01
18:46:09,472[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
40960 / 1281167 (3%)] epoch=0.0000 dt=0.6497 running_loss=0.0003
batch_loss=0.0087 acc=0.0016 batch_acc=0.0664[0m
[36m2022-11-01
18:46:38,547[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
46080 / 1281167 (4%)] epoch=0.0000 dt=0.7652 running_loss=0.0003
batch_loss=0.0088 acc=0.0018 batch_acc=0.0742[0m
[36m2022-11-01
18:47:08,492[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
51200 / 1281167 (4%)] epoch=0.0000 dt=0.7895 running_loss=0.0004
batch_loss=0.0090 acc=0.0022 batch_acc=0.0898[0m
[36m2022-11-01
18:47:40,341[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
56320 / 1281167 (4%)] epoch=0.0000 dt=0.6119 running_loss=0.0004
batch_loss=0.0088 acc=0.0025 batch_acc=0.0781[0m
[36m2022-11-01
18:48:11,392[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
61440 / 1281167 (5%)] epoch=0.0000 dt=0.5762 running_loss=0.0004
batch_loss=0.0085 acc=0.0029 batch_acc=0.0957[0m
[36m2022-11-01
18:48:40,259[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
66560 / 1281167 (5%)] epoch=0.0000 dt=0.6858 running_loss=0.0005
batch_loss=0.0085 acc=0.0033 batch_acc=0.1250[0m
[36m2022-11-01
18:49:26,019[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
71680 / 1281167 (6%)] epoch=0.0000 dt=0.6214 running_loss=0.0005
batch_loss=0.0085 acc=0.0037 batch_acc=0.1094[0m
[36m2022-11-01
18:49:57,591[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
76800 / 1281167 (6%)] epoch=0.0000 dt=0.6989 running_loss=0.0005
batch_loss=0.0086 acc=0.0041 batch_acc=0.0996[0m
[36m2022-11-01
18:50:30,173[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
81920 / 1281167 (6%)] epoch=0.0000 dt=0.9074 running_loss=0.0006
batch_loss=0.0086 acc=0.0045 batch_acc=0.0801[0m
[36m2022-11-01
18:51:03,423[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
```

87040 / 1281167 (7%)] epoch=0.0000 dt=0.5844 running_loss=0.0006
batch_loss=0.0083 acc=0.0050 batch_acc=0.1152[0m
[36m2022-11-01
18:51:37,800[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
92160 / 1281167 (7%)] epoch=0.0000 dt=0.7015 running_loss=0.0006
batch_loss=0.0084 acc=0.0054 batch_acc=0.1250[0m
[36m2022-11-01
18:52:11,099[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
97280 / 1281167 (8%)] epoch=0.0000 dt=0.5480 running_loss=0.0007
batch_loss=0.0084 acc=0.0059 batch_acc=0.1289[0m
[36m2022-11-01
18:52:46,464[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
102400 / 1281167 (8%)] epoch=0.0000 dt=0.6804 running_loss=0.0007
batch_loss=0.0082 acc=0.0063 batch_acc=0.1348[0m
[36m2022-11-01
18:53:15,315[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
107520 / 1281167 (8%)] epoch=0.0000 dt=0.5768 running_loss=0.0007
batch_loss=0.0080 acc=0.0068 batch_acc=0.1191[0m
[36m2022-11-01
18:53:43,585[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
112640 / 1281167 (9%)] epoch=0.0000 dt=0.6455 running_loss=0.0008
batch_loss=0.0079 acc=0.0073 batch_acc=0.1289[0m
[36m2022-11-01
18:54:13,286[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
117760 / 1281167 (9%)] epoch=0.0000 dt=0.7079 running_loss=0.0008
batch_loss=0.0081 acc=0.0078 batch_acc=0.1309[0m
[36m2022-11-01
18:54:42,195[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
122880 / 1281167 (10%)] epoch=0.0000 dt=0.8820 running_loss=0.0008
batch_loss=0.0079 acc=0.0084 batch_acc=0.1426[0m
[36m2022-11-01
18:55:18,540[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
128000 / 1281167 (10%)] epoch=0.0000 dt=0.6968 running_loss=0.0009
batch_loss=0.0081 acc=0.0089 batch_acc=0.1328[0m
[36m2022-11-01
18:55:46,297[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
133120 / 1281167 (10%)] epoch=0.0000 dt=0.5869 running_loss=0.0009
batch_loss=0.0080 acc=0.0095 batch_acc=0.1562[0m
[36m2022-11-01
18:56:19,046[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
138240 / 1281167 (11%)] epoch=0.0000 dt=0.7094 running_loss=0.0009
batch_loss=0.0080 acc=0.0100 batch_acc=0.1445[0m
[36m2022-11-01
18:56:50,962[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
143360 / 1281167 (11%)] epoch=0.0000 dt=0.9309 running_loss=0.0010
batch_loss=0.0081 acc=0.0106 batch_acc=0.1133[0m
[36m2022-11-01
18:57:19,155[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
148480 / 1281167 (12%)] epoch=0.0000 dt=0.7445 running_loss=0.0010
batch_loss=0.0079 acc=0.0112 batch_acc=0.1445[0m
[36m2022-11-01
18:57:52,957[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
153600 / 1281167 (12%)] epoch=0.0000 dt=0.7248 running_loss=0.0010
batch_loss=0.0079 acc=0.0117 batch_acc=0.1230[0m

```
[36m2022-11-01
18:58:26,024[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
158720 / 1281167 (12%)] epoch=0.0000 dt=0.7228 running_loss=0.0011
batch_loss=0.0079 acc=0.0123 batch_acc=0.1191[0m
[36m2022-11-01
18:58:56,783[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
163840 / 1281167 (13%)] epoch=0.0000 dt=0.7139 running_loss=0.0011
batch_loss=0.0080 acc=0.0130 batch_acc=0.1484[0m
[36m2022-11-01
18:59:25,941[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
168960 / 1281167 (13%)] epoch=0.0000 dt=0.5830 running_loss=0.0011
batch_loss=0.0080 acc=0.0136 batch_acc=0.1426[0m
[36m2022-11-01
18:59:58,149[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
174080 / 1281167 (14%)] epoch=0.0000 dt=0.8001 running_loss=0.0012
batch_loss=0.0076 acc=0.0142 batch_acc=0.1855[0m
[36m2022-11-01
19:00:28,896[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
179200 / 1281167 (14%)] epoch=0.0000 dt=0.6020 running_loss=0.0012
batch_loss=0.0076 acc=0.0148 batch_acc=0.1719[0m
[36m2022-11-01
19:00:59,381[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
184320 / 1281167 (14%)] epoch=0.0000 dt=0.5599 running_loss=0.0012
batch_loss=0.0079 acc=0.0155 batch_acc=0.1465[0m
[36m2022-11-01
19:01:27,167[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
189440 / 1281167 (15%)] epoch=0.0000 dt=0.7433 running_loss=0.0013
batch_loss=0.0078 acc=0.0162 batch_acc=0.1875[0m
[36m2022-11-01
19:01:59,935[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
194560 / 1281167 (15%)] epoch=0.0000 dt=0.7949 running_loss=0.0013
batch_loss=0.0076 acc=0.0168 batch_acc=0.1738[0m
[36m2022-11-01
19:02:28,256[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
199680 / 1281167 (16%)] epoch=0.0000 dt=0.6919 running_loss=0.0013
batch_loss=0.0080 acc=0.0174 batch_acc=0.1387[0m
[36m2022-11-01
19:02:57,946[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
204800 / 1281167 (16%)] epoch=0.0000 dt=0.5862 running_loss=0.0014
batch_loss=0.0078 acc=0.0181 batch_acc=0.1758[0m
[36m2022-11-01
19:03:27,797[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
209920 / 1281167 (16%)] epoch=0.0000 dt=0.6550 running_loss=0.0014
batch_loss=0.0077 acc=0.0188 batch_acc=0.1777[0m
[36m2022-11-01
19:04:01,135[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
215040 / 1281167 (17%)] epoch=0.0000 dt=0.6392 running_loss=0.0014
batch_loss=0.0078 acc=0.0194 batch_acc=0.1621[0m
[36m2022-11-01
19:04:33,303[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
220160 / 1281167 (17%)] epoch=0.0000 dt=0.8053 running_loss=0.0015
batch_loss=0.0076 acc=0.0201 batch_acc=0.1738[0m
[36m2022-11-01
19:05:03,174[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/1:
```

```
225280 / 1281167 (18%)] epoch=0.0000 dt=0.6149 running_loss=0.0015
batch_loss=0.0079 acc=0.0208 batch_acc=0.1875[0m
[36m2022-11-01
19:05:31,968[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
230400 / 1281167 (18%)] epoch=0.0000 dt=0.6193 running_loss=0.0015
batch_loss=0.0077 acc=0.0216 batch_acc=0.1973[0m
[36m2022-11-01
19:06:13,020[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
235520 / 1281167 (18%)] epoch=0.0000 dt=1.6355 running_loss=0.0015
batch_loss=0.0079 acc=0.0222 batch_acc=0.1855[0m
[36m2022-11-01
19:06:50,147[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
240640 / 1281167 (19%)] epoch=0.0000 dt=1.6329 running_loss=0.0016
batch_loss=0.0078 acc=0.0229 batch_acc=0.1816[0m
[36m2022-11-01
19:07:27,704[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/1:
245760 / 1281167 (19%)] epoch=0.0000 dt=1.4905 running_loss=0.0016
batch_loss=0.0075 acc=0.0237 batch_acc=0.1934[0m
```


2022-11-01-184210-thetagpu06_ngpu1_ncpu256.log

```
{
  "data": {
    "handler": "ilsvrc_dataset",
    "batch_size": 512,
    "train_filelist":
"/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
    "test_filelist":
"/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
    "shuffle_buffer": 200000,
    "reshuffle_each_iteration": true,
    "num_parallel_readers": 128,
    "prefetch_buffer_size": 8,
    "crop_image_size": [
      256,
      256
    ],
    "num_classes": 1000,
    "num_channels": 3
  }
}
num labels: 1000
build dataset /grand/projects/ALCFAITP/ilsvrc_train_filelist.txt
[36m2022-11-01
18:42:27,940[0m][34mai4sci.ilsvrc_dataset[0m][32mINFO[0m - input
filelist contains 1281167 files, estimated batches per rank 2502[0m
build dataset /grand/projects/ALCFAITP/ilsvrc_val_filelist.txt
[36m2022-11-01
18:42:31,884[0m][34mai4sci.ilsvrc_dataset[0m][32mINFO[0m - input
filelist contains 50000 files, estimated batches per rank 97[0m
```

2022-11-01-225933-thetagpu06_ngpu1_ncpu256

```
{
  "data": {
    "handler": "ilsvrc_dataset",
    "batch_size": 512,
    "train_filelist":
"/grand/projects/ALCFAITP/ilsvrc_train_filelist.txt",
    "test_filelist":
"/grand/projects/ALCFAITP/ilsvrc_val_filelist.txt",
    "shuffle_buffer": 200000,
    "reshuffle_each_iteration": true,
    "num_parallel_readers": 128,
    "prefetch_buffer_size": 8,
    "crop_image_size": [
      256,
      256
    ],
    "num_classes": 1000,
    "num_channels": 3
  }
}
num labels: 1000
build dataset /grand/projects/ALCFAITP/ilsvrc_train_filelist.txt
[36m2022-11-01
22:59:59,175[0m][34mai4sci.ilsvrc_dataset[0m][32mINFO[0m - input
filelist contains 1281167 files, estimated batches per rank 2502[0m
build dataset /grand/projects/ALCFAITP/ilsvrc_val_filelist.txt
[36m2022-11-01
23:00:01,382[0m][34mai4sci.ilsvrc_dataset[0m][32mINFO[0m - input
filelist contains 50000 files, estimated batches per rank 97[0m
thetagpu06:1527708:1527835 [0] NCCL INFO Bootstrap : Using
enp226s0:10.230.2.194<0>
thetagpu06:1527708:1527835 [0] NCCL INFO NET/Plugin : No plugin found
(libnccl-net.so), using internal implementation
thetagpu06:1527708:1527835 [0] NCCL INFO NET/IB : Using [0]mlx5_0:1/IB
[1]mlx5_1:1/IB [2]mlx5_2:1/IB [3]mlx5_3:1/IB [4]mlx5_4:1/IB
[5]mlx5_5:1/RoCE [6]mlx5_6:1/IB [7]mlx5_7:1/IB [8]mlx5_8:1/IB
[9]mlx5_9:1/IB [10]mlx5_10:1/IB [11]mlx5_11:1/RoCE [RO]; OOB
enp226s0:10.230.2.194<0>
thetagpu06:1527708:1527835 [0] NCCL INFO Using network IB
NCCL version 2.12.12+cuda11.0
thetagpu06:1527708:1527835 [0] NCCL INFO Setting affinity for GPU 0 to
010000,00000000,00000000,00000000,00010000,00000000
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 00/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 01/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 02/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 03/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 04/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 05/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 06/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 07/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 08/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 09/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 10/32 : 0
```

```
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 11/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 12/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 13/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 14/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 15/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 16/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 17/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 18/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 19/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 20/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 21/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 22/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 23/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 24/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 25/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 26/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 27/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 28/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 29/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 30/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Channel 31/32 : 0
thetagpu06:1527708:1527835 [0] NCCL INFO Trees [0] -1/-1/-1->0->-1 [1] -
1/-1/-1->0->-1 [2] -1/-1/-1->0->-1 [3] -1/-1/-1->0->-1 [4] -1/-1/-1->0->-
1 [5] -1/-1/-1->0->-1 [6] -1/-1/-1->0->-1 [7] -1/-1/-1->0->-1 [8] -1/-1/-
1->0->-1 [9] -1/-1/-1->0->-1 [10] -1/-1/-1->0->-1 [11] -1/-1/-1->0->-1
[12] -1/-1/-1->0->-1 [13] -1/-1/-1->0->-1 [14] -1/-1/-1->0->-1 [15] -1/-
1/-1->0->-1 [16] -1/-1/-1->0->-1 [17] -1/-1/-1->0->-1 [18] -1/-1/-1->0->-
1 [19] -1/-1/-1->0->-1 [20] -1/-1/-1->0->-1 [21] -1/-1/-1->0->-1 [22] -
1/-1/-1->0->-1 [23] -1/-1/-1->0->-1 [24] -1/-1/-1->0->-1 [25] -1/-1/-1-
>0->-1 [26] -1/-1/-1->0->-1 [27] -1/-1/-1->0->-1 [28] -1/-1/-1->0->-1
[29] -1/-1/-1->0->-1 [30] -1/-1/-1->0->-1 [31] -1/-1/-1->0->-1
thetagpu06:1527708:1527835 [0] NCCL INFO Connected all rings
thetagpu06:1527708:1527835 [0] NCCL INFO Connected all trees
thetagpu06:1527708:1527835 [0] NCCL INFO 32 coll channels, 32 p2p
channels, 32 p2p channels per peer
thetagpu06:1527708:1527835 [0] NCCL INFO comm 0x7f880afff6d0 rank 0
nranks 1 cudaDev 0 busId f000 - Init COMPLETE
[36m2022-11-01
23:00:41,917[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5: 0 /
1281167 (0%)] epoch=0.0000 dt=37.8666 running_loss=0.0000
batch_loss=0.0151 acc=0.0000 batch_acc=0.0039[0m
[36m2022-11-01
23:00:56,684[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
5120 / 1281167 (0%)] epoch=0.0000 dt=0.7848 running_loss=0.0000
batch_loss=0.0095 acc=0.0001 batch_acc=0.0293[0m
[36m2022-11-01
23:01:17,685[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
10240 / 1281167 (1%)] epoch=0.0000 dt=0.6010 running_loss=0.0001
batch_loss=0.0095 acc=0.0002 batch_acc=0.0430[0m
[36m2022-11-01
23:01:41,301[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
15360 / 1281167 (1%)] epoch=0.0000 dt=0.7152 running_loss=0.0001
batch_loss=0.0093 acc=0.0004 batch_acc=0.0449[0m
```

[36m2022-11-01
23:02:04,130[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
20480 / 1281167 (2%)] epoch=0.0000 dt=0.7532 running_loss=0.0002
batch_loss=0.0095 acc=0.0006 batch_acc=0.0410[0m
[36m2022-11-01
23:02:27,030[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
25600 / 1281167 (2%)] epoch=0.0000 dt=0.6965 running_loss=0.0002
batch_loss=0.0091 acc=0.0008 batch_acc=0.0391[0m
[36m2022-11-01
23:03:02,037[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
30720 / 1281167 (2%)] epoch=0.0000 dt=1.6816 running_loss=0.0002
batch_loss=0.0090 acc=0.0010 batch_acc=0.0703[0m
[36m2022-11-01
23:03:24,701[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
35840 / 1281167 (3%)] epoch=0.0000 dt=0.7752 running_loss=0.0003
batch_loss=0.0091 acc=0.0012 batch_acc=0.0605[0m
[36m2022-11-01
23:03:48,793[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
40960 / 1281167 (3%)] epoch=0.0000 dt=0.6317 running_loss=0.0003
batch_loss=0.0089 acc=0.0015 batch_acc=0.0859[0m
[36m2022-11-01
23:04:12,383[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
46080 / 1281167 (4%)] epoch=0.0000 dt=0.6728 running_loss=0.0003
batch_loss=0.0090 acc=0.0018 batch_acc=0.0684[0m
[36m2022-11-01
23:04:35,957[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
51200 / 1281167 (4%)] epoch=0.0000 dt=0.7066 running_loss=0.0004
batch_loss=0.0088 acc=0.0021 batch_acc=0.0840[0m
[36m2022-11-01
23:04:59,412[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
56320 / 1281167 (4%)] epoch=0.0000 dt=0.7741 running_loss=0.0004
batch_loss=0.0090 acc=0.0024 batch_acc=0.0625[0m
[36m2022-11-01
23:05:23,152[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
61440 / 1281167 (5%)] epoch=0.0000 dt=0.7233 running_loss=0.0004
batch_loss=0.0090 acc=0.0027 batch_acc=0.0645[0m
[36m2022-11-01
23:05:53,969[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
66560 / 1281167 (5%)] epoch=0.0000 dt=1.0540 running_loss=0.0005
batch_loss=0.0085 acc=0.0031 batch_acc=0.1172[0m
[36m2022-11-01
23:06:18,065[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
71680 / 1281167 (6%)] epoch=0.0000 dt=0.6233 running_loss=0.0005
batch_loss=0.0085 acc=0.0035 batch_acc=0.0879[0m
[36m2022-11-01
23:06:42,057[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
76800 / 1281167 (6%)] epoch=0.0000 dt=0.7227 running_loss=0.0006
batch_loss=0.0083 acc=0.0039 batch_acc=0.1230[0m
[36m2022-11-01
23:07:09,798[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:
81920 / 1281167 (6%)] epoch=0.0000 dt=0.6721 running_loss=0.0006
batch_loss=0.0086 acc=0.0043 batch_acc=0.1113[0m
[36m2022-11-01
23:07:35,110[0m][34mai4sci.trainer[0m][32mINFO[0m - [0] [0/5:

87040 / 1281167 (7%)] epoch=0.0000 dt=0.7298 running_loss=0.0006
batch_loss=0.0085 acc=0.0047 batch_acc=0.1152[0m
[36m2022-11-01
23:07:58,440[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
92160 / 1281167 (7%)] epoch=0.0000 dt=0.6893 running_loss=0.0007
batch_loss=0.0084 acc=0.0051 batch_acc=0.1016[0m
[36m2022-11-01
23:08:30,941[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
97280 / 1281167 (8%)] epoch=0.0000 dt=0.8261 running_loss=0.0007
batch_loss=0.0084 acc=0.0055 batch_acc=0.1113[0m
[36m2022-11-01
23:08:54,423[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
102400 / 1281167 (8%)] epoch=0.0000 dt=0.7907 running_loss=0.0007
batch_loss=0.0081 acc=0.0060 batch_acc=0.1270[0m
[36m2022-11-01
23:09:22,263[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
107520 / 1281167 (8%)] epoch=0.0000 dt=0.6856 running_loss=0.0008
batch_loss=0.0079 acc=0.0065 batch_acc=0.1465[0m
[36m2022-11-01
23:09:46,562[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
112640 / 1281167 (9%)] epoch=0.0000 dt=0.8142 running_loss=0.0008
batch_loss=0.0083 acc=0.0070 batch_acc=0.1270[0m
[36m2022-11-01
23:10:10,363[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
117760 / 1281167 (9%)] epoch=0.0000 dt=1.0331 running_loss=0.0008
batch_loss=0.0078 acc=0.0075 batch_acc=0.1816[0m
[36m2022-11-01
23:10:34,879[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
122880 / 1281167 (10%)] epoch=0.0000 dt=0.9838 running_loss=0.0009
batch_loss=0.0080 acc=0.0080 batch_acc=0.1406[0m
[36m2022-11-01
23:11:00,533[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
128000 / 1281167 (10%)] epoch=0.0000 dt=0.9275 running_loss=0.0009
batch_loss=0.0082 acc=0.0085 batch_acc=0.1289[0m
[36m2022-11-01
23:11:26,533[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
133120 / 1281167 (10%)] epoch=0.0000 dt=1.6944 running_loss=0.0009
batch_loss=0.0081 acc=0.0090 batch_acc=0.1504[0m
[36m2022-11-01
23:11:55,199[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
138240 / 1281167 (11%)] epoch=0.0000 dt=0.9835 running_loss=0.0010
batch_loss=0.0082 acc=0.0096 batch_acc=0.1309[0m
[36m2022-11-01
23:12:18,729[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
143360 / 1281167 (11%)] epoch=0.0000 dt=0.6404 running_loss=0.0010
batch_loss=0.0080 acc=0.0101 batch_acc=0.1426[0m
[36m2022-11-01
23:12:42,449[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
148480 / 1281167 (12%)] epoch=0.0000 dt=0.7136 running_loss=0.0010
batch_loss=0.0083 acc=0.0107 batch_acc=0.1348[0m
[36m2022-11-01
23:13:08,832[0m][34mai4sci.trainer[0m][32mINFO[0m] - [0] [0/5:
153600 / 1281167 (12%)] epoch=0.0000 dt=0.7825 running_loss=0.0010
batch_loss=0.0084 acc=0.0113 batch_acc=0.1387[0m

[36m2022-11-01
23:13:32,600[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
158720 / 1281167 (12%)] epoch=0.0000 dt=0.8286 running_loss=0.0011
batch_loss=0.0084 acc=0.0118 batch_acc=0.1230[0m
[36m2022-11-01
23:13:56,610[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
163840 / 1281167 (13%)] epoch=0.0000 dt=0.9678 running_loss=0.0011
batch_loss=0.0080 acc=0.0124 batch_acc=0.1484[0m
[36m2022-11-01
23:14:23,151[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
168960 / 1281167 (13%)] epoch=0.0000 dt=0.6450 running_loss=0.0011
batch_loss=0.0077 acc=0.0131 batch_acc=0.1621[0m
[36m2022-11-01
23:14:54,279[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
174080 / 1281167 (14%)] epoch=0.0000 dt=0.6394 running_loss=0.0012
batch_loss=0.0079 acc=0.0137 batch_acc=0.1680[0m
[36m2022-11-01
23:15:18,828[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
179200 / 1281167 (14%)] epoch=0.0000 dt=0.6705 running_loss=0.0012
batch_loss=0.0081 acc=0.0143 batch_acc=0.1348[0m
[36m2022-11-01
23:15:42,609[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
184320 / 1281167 (14%)] epoch=0.0000 dt=0.6772 running_loss=0.0012
batch_loss=0.0076 acc=0.0149 batch_acc=0.1426[0m
[36m2022-11-01
23:16:09,013[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
189440 / 1281167 (15%)] epoch=0.0000 dt=0.6017 running_loss=0.0013
batch_loss=0.0080 acc=0.0155 batch_acc=0.1523[0m
[36m2022-11-01
23:16:33,861[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
194560 / 1281167 (15%)] epoch=0.0000 dt=0.9773 running_loss=0.0013
batch_loss=0.0081 acc=0.0161 batch_acc=0.1504[0m
[36m2022-11-01
23:16:59,377[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
199680 / 1281167 (16%)] epoch=0.0000 dt=0.8887 running_loss=0.0013
batch_loss=0.0079 acc=0.0168 batch_acc=0.1738[0m
[36m2022-11-01
23:17:23,695[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
204800 / 1281167 (16%)] epoch=0.0000 dt=1.1651 running_loss=0.0014
batch_loss=0.0081 acc=0.0175 batch_acc=0.1660[0m
[36m2022-11-01
23:17:55,903[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
209920 / 1281167 (16%)] epoch=0.0000 dt=0.6303 running_loss=0.0014
batch_loss=0.0080 acc=0.0182 batch_acc=0.1543[0m
[36m2022-11-01
23:18:20,857[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
215040 / 1281167 (17%)] epoch=0.0000 dt=0.5843 running_loss=0.0014
batch_loss=0.0081 acc=0.0188 batch_acc=0.1523[0m
[36m2022-11-01
23:18:46,716[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
220160 / 1281167 (17%)] epoch=0.0000 dt=0.8611 running_loss=0.0015
batch_loss=0.0080 acc=0.0195 batch_acc=0.1523[0m
[36m2022-11-01
23:19:10,647[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:

225280 / 1281167 (18%)] epoch=0.0000 dt=0.6022 running_loss=0.0015
batch_loss=0.0077 acc=0.0202 batch_acc=0.1602[0m
[0[36m2022-11-01
23:19:37,339[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
230400 / 1281167 (18%)] epoch=0.0000 dt=0.8068 running_loss=0.0015
batch_loss=0.0075 acc=0.0209 batch_acc=0.1992[0m
[0[36m2022-11-01
23:20:01,158[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
235520 / 1281167 (18%)] epoch=0.0000 dt=0.6996 running_loss=0.0016
batch_loss=0.0079 acc=0.0216 batch_acc=0.1504[0m
[0[36m2022-11-01
23:20:33,017[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
240640 / 1281167 (19%)] epoch=0.0000 dt=0.6435 running_loss=0.0016
batch_loss=0.0081 acc=0.0223 batch_acc=0.1562[0m
[0[36m2022-11-01
23:20:59,565[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
245760 / 1281167 (19%)] epoch=0.0000 dt=0.9551 running_loss=0.0016
batch_loss=0.0078 acc=0.0230 batch_acc=0.1699[0m
[0[36m2022-11-01
23:21:23,950[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
250880 / 1281167 (20%)] epoch=0.0000 dt=0.6634 running_loss=0.0016
batch_loss=0.0079 acc=0.0238 batch_acc=0.1680[0m
[0[36m2022-11-01
23:21:49,697[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
256000 / 1281167 (20%)] epoch=0.0000 dt=0.7473 running_loss=0.0017
batch_loss=0.0080 acc=0.0245 batch_acc=0.1680[0m
[0[36m2022-11-01
23:22:15,719[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
261120 / 1281167 (20%)] epoch=0.0000 dt=0.7452 running_loss=0.0017
batch_loss=0.0079 acc=0.0252 batch_acc=0.1562[0m
[0[36m2022-11-01
23:22:39,306[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
266240 / 1281167 (21%)] epoch=0.0000 dt=0.8291 running_loss=0.0017
batch_loss=0.0078 acc=0.0260 batch_acc=0.1816[0m
[0[36m2022-11-01
23:23:10,853[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
271360 / 1281167 (21%)] epoch=0.0000 dt=0.7267 running_loss=0.0018
batch_loss=0.0077 acc=0.0267 batch_acc=0.1953[0m
[0[36m2022-11-01
23:23:36,541[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
276480 / 1281167 (22%)] epoch=0.0000 dt=1.0576 running_loss=0.0018
batch_loss=0.0075 acc=0.0275 batch_acc=0.1797[0m
[0[36m2022-11-01
23:24:16,968[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
281600 / 1281167 (22%)] epoch=0.0000 dt=0.6180 running_loss=0.0018
batch_loss=0.0073 acc=0.0283 batch_acc=0.2051[0m
[0[36m2022-11-01
23:25:12,228[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
286720 / 1281167 (22%)] epoch=0.0000 dt=0.6734 running_loss=0.0019
batch_loss=0.0073 acc=0.0291 batch_acc=0.2070[0m
[0[36m2022-11-01
23:25:38,558[0m] [0[34mai4sci.trainer[0m] [0[32mINFO[0m] - [0] [0/5:
291840 / 1281167 (23%)] epoch=0.0000 dt=0.6276 running_loss=0.0019
batch_loss=0.0071 acc=0.0299 batch_acc=0.2344[0m

```
[36m2022-11-01
23:26:11,529[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
296960 / 1281167 (23%)] epoch=0.0000 dt=8.5706 running_loss=0.0019
batch_loss=0.0075 acc=0.0307 batch_acc=0.2070[0m
[36m2022-11-01
23:26:35,076[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
302080 / 1281167 (24%)] epoch=0.0000 dt=0.6785 running_loss=0.0019
batch_loss=0.0075 acc=0.0315 batch_acc=0.2031[0m
[36m2022-11-01
23:27:02,577[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
307200 / 1281167 (24%)] epoch=0.0000 dt=0.6337 running_loss=0.0020
batch_loss=0.0073 acc=0.0323 batch_acc=0.2168[0m
[36m2022-11-01
23:27:27,867[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
312320 / 1281167 (24%)] epoch=0.0000 dt=0.6733 running_loss=0.0020
batch_loss=0.0077 acc=0.0331 batch_acc=0.2012[0m
[36m2022-11-01
23:27:51,985[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
317440 / 1281167 (25%)] epoch=0.0000 dt=0.7596 running_loss=0.0020
batch_loss=0.0071 acc=0.0339 batch_acc=0.2422[0m
[36m2022-11-01
23:28:17,810[0m] [34mai4sci.trainer[0m] [32mINFO[0m] - [0] [0/5:
322560 / 1281167 (25%)] epoch=0.0000 dt=0.7298 running_loss=0.0021
batch_loss=0.0073 acc=0.0348 batch_acc=0.2227[0m
```